


Analytical scenario: 60 million parts + 1 million operating manuals



Explicit assumptions

Parts: $60\text{M} \times 500 \text{ tokens} = 30\text{B}$. Manuals: $1\text{M} \times 10,000 \text{ tokens} = 10\text{B}$.

Sensitivity: 100–2,000 tokens/part gives 6,000–120,000 prompt equivalents.

At 500 tokens/part, a 1M-token prompt holds at most 2,000 parts (0.0033%).

Capacity illustration, not measured REA throughput, accuracy, or exhaustive coverage.